

# Execution time speedup over serial kernel scheduling

Block size:


● 1D=32, 2D=8

✕ 1D=128, 2D=8


◆ 1D=256, 2D=16

+ 1D=1024, 2D=32


Vector Squares




ML Ensemble



HITS



Images



Input number of elements